

# **Logistic Regression**

Analysed by: Pavan Eleti

#### Introduction:

Loantap is a leading financial technology company based in India, specializing in providing flexible and innovative loan products to individuals and businesses. With a focus on customercentric solutions, Loantap leverages technology to offer hassle-free borrowing experiences, including personal loans, salary advances, and flexible EMI options. Their commitment to transparency, speed, and convenience has established them as a trusted partner for borrowers seeking efficient financial solutions.

- LoanTap is at the forefront of offering tailored financial solutions to millennials.
- Their innovative approach seeks to harness data science for refining their credit underwriting process.
- The focus here is the Personal Loan segment. A deep dive into the dataset can reveal patterns in borrower behavior and creditworthiness.
- Analyzing this dataset can provide crucial insights into the financial behaviors, spending habits, and potential risk associated with each borrower.
- The insights gained can optimize loan disbursal, balancing customer outreach with risk management.

#### Our Task:

As a data scientist at LoanTap, you are tasked with analyzing the dataset to
determine the creditworthiness of potential borrowers. Your ultimate
objective is to build a logistic regression model, evaluate its performance,
and provide actionable insights for the underwriting process.

### Features of the dataset:

• Column Profiling:

*	1 1.
Feature	Description
loan_amnt	The listed amount of the loan applied for by the borrower. If at some point in time, the credit department
term	The number of payments on the loan. Values are in months and can be either 36 or 60
int_rate	Interest Rate on the loan
installment	The monthly payment owed by the borrower if the loan originates
grade	LoanTap assigned loan grade
sub_grade	LoanTap assigned loan subgrade
emp_title	The job title supplied by the Borrower when applying for the loan
emp_length	Employment length in years. Possible values are between 0 and 10 where 0 means less than one year ar
home_ownership	The home ownership status provided by the borrower during registration or obtained from the credit repo
annual_inc	The self-reported annual income provided by the borrower during registration
verification_status	Indicates if income was verified by LoanTap, not verified, or if the income source was verified
issue_d	The month which the loan was funded
loan_status	Current status of the loan - Target Variable
purpose	A category provided by the borrower for the loan request
title	The loan title provided by the borrower
dti	A ratio calculated using the borrower's total monthly debt payments on the total debt obligations, exclud
earliest_cr_line	The month the borrower's earliest reported credit line was opened
open_acc	The number of open credit lines in the borrower's credit file
pub_rec	Number of derogatory public records
revol_bal	Total credit revolving balance
revol_util	Revolving line utilization rate, or the amount of credit the borrower is using relative to all available revolvi
total_acc	The total number of credit lines currently in the borrower's credit file
initial_list_status	The initial listing status of the loan
application_type	Indicates whether the loan is an individual application or a joint application with two co-borrowers
mort_acc	Number of mortgage accounts
pub_rec_bankruptcies	Number of public record bankruptcies
Address	Address of the individual

# 

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import ttest ind,chi2 contingency
from sklearn.linear_model import LogisticRegression
from sklearn.model selection import train test split, KFold, cross val score
from sklearn.preprocessing import MinMaxScaler, LabelEncoder, StandardScaler
from sklearn.metrics import (
    accuracy_score, confusion_matrix, classification_report,
    roc_auc_score, roc_curve, auc, precision_recall_curve, average_precision_scor
    ConfusionMatrixDisplay, RocCurveDisplay, f1_score, recall_score, precision_score
)
from statsmodels.stats.outliers_influence import variance_inflation_factor
from imblearn.over_sampling import SMOTE
import warnings
warnings.filterwarnings("ignore")
!qdown 1sutK5BbP4CEMbnB9hq00K75uFW754scz
→ Downloading...
     From: <a href="https://drive.google.com/uc?id=1sutK5BbP4CEMbnB9hq00K75uFW754scz">https://drive.google.com/uc?id=1sutK5BbP4CEMbnB9hq00K75uFW754scz</a>
    To: /content/loantap.csv
     100% 100M/100M [00:01<00:00, 53.6MB/s]
lt_data =pd.read_csv('loantap-data.csv')
df = lt_data.copy()
df.head()
```

loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_l			
10000.0	36 months	11.44	329.48	В	В4	Marketing	10-			
8000.0	36 months	11.99	265.68	В	B5	Credit analyst	۷			
15600.0	36 months	10.49	506.97	В	В3	Statistician	<			
7200.0	36 months	6.49	220.65	А	A2	Client Advocate	6			
24375.0	60 months	17.27	609.33	С	C5	Destiny Management Inc.	ξ			
5 rows × 27 columns										
	10000.0 8000.0 15600.0 7200.0	10000.0 36 months  8000.0 36 months  15600.0 36 months  7200.0 36 months  24375.0 60 months	10000.0 36 11.44  8000.0 36 11.99  15600.0 36 10.49  7200.0 36 months 6.49  24375.0 60 months 17.27	10000.0 36 11.44 329.48  8000.0 36 11.99 265.68  15600.0 36 10.49 506.97  7200.0 36 months 6.49 220.65  24375.0 60 months 17.27 609.33	10000.0 36 months 11.44 329.48 B  8000.0 36 months 11.99 265.68 B  15600.0 36 months 10.49 506.97 B  7200.0 36 months 6.49 220.65 A  24375.0 60 months 17.27 609.33 C	10000.0       36 months       11.44       329.48       B       B4         8000.0       36 months       11.99       265.68       B       B5         15600.0       36 months       10.49       506.97       B       B3         7200.0       36 months       6.49       220.65       A       A2         24375.0       60 months       17.27       609.33       C       C5	10000.0       36 months       11.44       329.48       B       B4       Marketing         8000.0       36 months       11.99       265.68       B       B5       Credit analyst         15600.0       36 months       10.49       506.97       B       B3       Statistician         7200.0       36 months       6.49       220.65       A       A2       Client Advocate         24375.0       60 months       17.27       609.33       C       C5       Management Inc.			

pd.set\_option('display.max\_columns', None)

# Exploration of data :

df.shape

**→** (396030, 27)

df.info()

<<rp><class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):

#	Column	Non-Null Count	Dtype
0	loan_amnt	396030 non-null	float64
1	term	396030 non-null	object
2	int_rate	396030 non-null	float64
3	installment	396030 non-null	float64
4	grade	396030 non-null	object
5	sub_grade	396030 non-null	object
6	emp_title	373103 non-null	object
7	emp_length	377729 non-null	object
8	home_ownership	396030 non-null	object
9	annual_inc	396030 non-null	float64
10	verification_status	396030 non-null	object
11	issue_d	396030 non-null	object
12	loan_status	396030 non-null	object
13	purpose	396030 non-null	object
14	title	394274 non-null	object

```
dti
                         396030 non-null float64
15
                         396030 non-null object
16 earliest_cr_line
17
   open_acc
                         396030 non-null
                                          float64
                         396030 non-null
                                          float64
18
   pub rec
   revol_bal
19
                         396030 non-null float64
20 revol_util
                         395754 non-null float64
21
   total_acc
                         396030 non-null float64
22
   initial_list_status
                         396030 non-null object
23 application_type
                         396030 non-null object
24 mort_acc
                         358235 non-null
                                          float64
25 pub_rec_bankruptcies
                         395495 non-null float64
26 address
                         396030 non-null object
```

dtypes: float64(12), object(15)

memory usage: 81.6+ MB

#### df.columns

## Statistical Summary

#### df.describe().T

<b>→</b>		count	mean	std	min	25%	50%
	loan_amnt	396030.0	14113.888089	8357.441341	500.00	8000.00	12000.00
	int_rate	396030.0	13.639400	4.472157	5.32	10.49	13.33
	installment	396030.0	431.849698	250.727790	16.08	250.33	375.43
	annual_inc	396030.0	74203.175798	61637.621158	0.00	45000.00	64000.00
	dti	396030.0	17.379514	18.019092	0.00	11.28	16.91
	open_acc	396030.0	11.311153	5.137649	0.00	8.00	10.00
	pub_rec	396030.0	0.178191	0.530671	0.00	0.00	0.00
	revol_bal	396030.0	15844.539853	20591.836109	0.00	6025.00	11181.00
	revol_util	395754.0	53.791749	24.452193	0.00	35.80	54.80
	total_acc	396030.0	25.414744	11.886991	2.00	17.00	24.00
	mort_acc	358235.0	1.813991	2.147930	0.00	0.00	1.00
							• • •

df.describe(include='object').T

_		_
•	₹	_
_		$\overline{}$
1		_

	count	unique	top	freq
term	396030	2	36 months	302005
grade	396030	7	В	116018
sub_grade	396030	35	B3	26655
emp_title	373103	173105	Teacher	4389
emp_length	377729	11	10+ years	126041
home_ownership	396030	6	MORTGAGE	198348
verification_status	396030	3	Verified	139563
issue_d	396030	115	Oct-2014	14846
loan_status	396030	2	Fully Paid	318357
purpose	396030	14	debt_consolidation	234507
title	394274	48816	Debt consolidation	152472
earliest_cr_line	396030	684	Oct-2000	3017
initial_list_status	396030	2	f	238066
application_type	396030	3	INDIVIDUAL	395319
address	396030	393700	USCGC Smith\r\nFPO AE 70466	8

# → ■■ Duplicate Detection

df[df.duplicated()]



loan\_amnt term int\_rate installment grade sub\_grade emp\_title emp\_leng



• The dataset does not contain any duplicates.

# ? Null Detection

df.isna().any()[df.isna().any()]

$\rightarrow$	emp_title	True
	emp_length	True
	title	True
	revol_util	True
	mort_acc	True

pub\_rec\_bankruptcies True
dtype: bool

df.isna().sum().sort\_values(ascending=False)

<b>→</b>	<pre>mort_acc emp_title</pre>	37795 22927
	emp_length	18301
	title	1756
	<pre>pub_rec_bankruptcies</pre>	535
	revol_util	276
	loan_amnt	0
	dti	0
	application_type	0
	initial_list_status	0
	total_acc	0
	revol_bal	0
	pub_rec	0
	open_acc	0
	earliest_cr_line	0
	purpose	0
	term	0
	loan_status	0
	issue_d	0
	verification_status	0
	annual_inc	0
	home_ownership	0
	sub_grade	0
	grade	0
	installment	0
	int_rate	0
	address	0
	dtype: int64	_
	• •	

def missing\_data(df):

total\_missing\_df = df.isnull().sum().sort\_values(ascending =False)
percent\_missing\_df = (df.isnull().sum()/df.isna().count()\*100).sort\_values(as
missing\_data\_df = pd.concat([total\_missing\_df, percent\_missing\_df], axis=1, k
return missing\_data\_df

missing\_pct = missing\_data(df)
missing\_pct[missing\_pct['Total']>0]

<b>→</b> ▼		Total	Percent
	mort_acc	37795	9.543469
	emp_title	22927	5.789208
	emp_length	18301	4.621115
	title	1756	0.443401
	pub_rec_bankruptcies	535	0.135091
	revol util	276	0.069692



#### Following columns has missing values

- 1. emp\_title has 5.78% missing values
- 2. emp\_length has 4.62% missing values
- 3. title has 0.44% missing values
- 4. revol\_until has 0.06% missing values
- 5. mort\_acc has 9.54% missing values
- 6. pub\_rec\_bankruptcies has 0.13% missing values

#### Action

 Since ML algorithm do not work on columns which has missing values so we need to impute these missing values.

```
plt.figure(figsize=(25,8))
plt.style.use('dark_background')
sns.heatmap(df.isnull().T,cmap='Purples')
plt.title('Visual Check of Nulls',fontsize=20,color='magenta')
plt.show()
```



```
\overline{\Rightarrow}
```

```
Total Unique Values in loan amnt column are :- 1397
Unique Values in loan_amnt column are :-
 [10000. 8000. 15600. ... 36275. 36475.
                                          725.]
Value_counts of loan_amnt column :-
 loan amnt
10000.0
           27668
12000.0
           21366
15000.0
          19903
20000.0
           18969
35000.0
          14576
36225.0
              1
950.0
              1
              1
37800.0
30050.0
              1
725.0
              1
Name: count, Length: 1397, dtype: int64
Total Unique Values in term column are :- 2
Unique Values in term column are :-
 [' 36 months' ' 60 months']
Value_counts of term column :-
 term
 36 months
             302005
 60 months
              94025
Name: count, dtype: int64
Total Unique Values in int rate column are :- 566
Unique Values in int rate column are :-
 [11.44 11.99 10.49 6.49 17.27 13.33 5.32 11.14 10.99 16.29 13.11 14.64
  9.17 12.29 6.62 8.39 21.98 7.9 6.97 6.99 15.61 11.36 13.35 12.12
  9.99 8.19 18.75 6.03 14.99 16.78 13.67 13.98 16.99 19.91 17.86 21.49
 12.99 18.54 7.89 17.1 18.25 11.67 6.24 8.18 12.35 14.16 17.56 18.55
 22.15 10.39 15.99 16.07 24.99 9.67 19.19 21.
                                                12.69 10.74 6.68 19.22
 11.49 16.55 19.97 24.7 13.49 18.24 16.49 25.78 25.83 18.64
            7.69 19.53 10.16 7.62 9.75 13.68 15.88 14.65 6.92 23.83
 15.22 15.31
 10.75 18.49 20.31 17.57 27.31 19.99 22.99 12.59 10.37 14.33 13.53 22.45
 24.5 17.99 9.16 12.49 11.55 17.76 28.99 23.1 20.49 22.7 10.15 6.89
 19.52 8.9 14.3
                   9.49 25.99 24.08 13.05 14.98 16.59 11.26 25.89 14.48
 21.99 23.99 5.99 14.47 11.53 8.67 8.59 10.64 23.28 25.44 9.71 16.2
 19.24 24.11 15.8 15.96 14.49 18.99 5.79 19.29 14.54 14.09 9.25 19.05
 17.77 18.92 20.75 10.65 18.85 10.59 12.85 11.39 13.65 13.06 7.12 20.99
 13.61 12.73 14.46 16.24 25.49 7.39 10.78 20.8
                                                7.88 15.95 12.39 21.18
 21.97 15.77
            6.39 10.
                        12.53 13.43 7.49 25.57 21.48 18.39 11.47
 15.68 19.04 14.31 24.24 5.42 23.43 19.47 6.54 23.32 17.58 14.72
                                                                  7.66
 9.76 13.23 13.48 12.42 9.8 11.71 14.27 21.15 22.95 8.49 17.74 15.59
 13.72 9.45
            7.29 15.1 11.86 19.72 14.35 11.22 15.62 15.81 12.41 28.67
 11.48 13.66 9.91 23.76 17.14 18.84 12.23 6.17 8.94 14.22 19.03 25.29
  8.99 9.88 15.58 27.49 8.07 22.47 19.2 13.44 22.4
                                                      12.79 18.2
                                                                  13.18
  7.24 14.84 5.93 15.28 13.85 25.28 8.
                                          9.62 12.05 15.7 20.2
                                                                  13.57
       7.4 25.8 12.68 11.83 7.37 11.11 14.85 16.
                                                      11.12 23.63
 21.67
 7 00 7 01 17 02 21 7 26 66 16 77 27 27 12 21 7 60 15 27 10 60 0 62
```

## Null Treatment:

```
df.loc[df['revol_util'].isna(), 'revol_util'] = 0.0
df.loc[df['mort_acc'].isna(),'mort_acc'] = 0.0
df.loc[df['pub_rec_bankruptcies'].isna(),'pub_rec_bankruptcies'] = 0.0
df.loc[df['emp_title'].isna(),'emp_title'] = 'No Employee Title'
df.loc[df['title'].isna(),'title'] = 'Unavailable'
df['emp_length'] = df['emp_length'].fillna('< 1 year')</pre>
df.isna().sum()
→ loan amnt
                              0
     term
                              0
     int rate
                              0
                              0
     installment
                              0
     grade
     sub grade
                              0
     emp_title
                              0
     emp_length
                              0
    home_ownership
                              0
    annual_inc
     verification_status
                              0
     issue d
                              0
     loan_status
                              0
    purpose
                              0
    title
                              0
    dti
                              0
     earliest_cr_line
                              0
                              0
     open_acc
     pub_rec
                              0
                              0
     revol_bal
     revol_util
                              0
     total_acc
                              0
     initial_list_status
                              0
     application_type
                              0
    mort_acc
                              0
                              0
     pub_rec_bankruptcies
                              0
     address
     dtype: int64
```

df.describe().T



	count	mean	std	min	25%	50%
loan_amnt	396030.0	14113.888089	8357.441341	500.00	8000.00	12000.00
int_rate	396030.0	13.639400	4.472157	5.32	10.49	13.33
installment	396030.0	431.849698	250.727790	16.08	250.33	375.43
annual_inc	396030.0	74203.175798	61637.621158	0.00	45000.00	64000.00
dti	396030.0	17.379514	18.019092	0.00	11.28	16.91
open_acc	396030.0	11.311153	5.137649	0.00	8.00	10.00
pub_rec	396030.0	0.178191	0.530671	0.00	0.00	0.00
revol_bal	396030.0	15844.539853	20591.836109	0.00	6025.00	11181.00
revol_util	396030.0	53.754260	24.484857	0.00	35.80	54.80
total_acc	396030.0	25.414744	11.886991	2.00	17.00	24.00
mort_acc	396030.0	1.640873	2.111249	0.00	0.00	1.00
		2 121 122		^ ^^		

df.describe(include='object').T

<b>→</b>		_			_
``		count	unique	top	freq
	term	396030	2	36 months	302005
	grade	396030	7	В	116018
	sub_grade	396030	35	В3	26655
	emp_title	396030	173106	No Employee Title	22927
	emp_length	396030	11	10+ years	126041
	home_ownership	396030	6	MORTGAGE	198348
	verification_status	396030	3	Verified	139563
	issue_d	396030	115	Oct-2014	14846
	loan_status	396030	2	Fully Paid	318357
	purpose	396030	14	debt_consolidation	234507
	title	396030	48817	Debt consolidation	152472
	earliest_cr_line	396030	684	Oct-2000	3017
	initial_list_status	396030	2	f	238066
	application_type	396030	3	INDIVIDUAL	395319
	address	396030	393700	USCGC Smith\r\nFPO AE 70466	8

## , 🌅 Feature Engineering

```
df['pub_rec'] = [1 if i > 1 else 0 for i in df['pub_rec']]
df['mort acc'] = [1 if i > 1 else 0 for i in <math>df['mort acc']]
df['pub rec bankruptcies'] = [1 if i > 1 else 0 for i in df['pub rec bankruptcies
df.sample()
\rightarrow
            loan_amnt
                        term int_rate installment grade sub_grade
                                                                           emp_title
                                                                        Superintendent
                           36
                                                                    C1
     60136
               35000.0
                                   12.29
                                              1167.36
                                                          C
                       months
                                                                          Maintenance
#Split issue_date into month and year
df[['issue_month', 'issue_year']] = df['issue_d'].str.split('-', expand=True)
df.drop(['issue d'], axis=1, inplace=True)
#Split er_cr_line date into month and year
df[['er cr line m', 'er cr line y']] = df['earliest cr line'].str.split('-', expa
df.drop(['earliest cr line'], axis=1, inplace=True)
df['address']
\rightarrow
                  0174 Michelle Gateway\r\nMendozaberg, OK 22690
    1
               1076 Carney Fort Apt. 347\r\nLoganmouth, SD 05113
    2
               87025 Mark Dale Apt. 269\r\nNew Sabrina, WV 05113
    3
                         823 Reid Ford\r\nDelacruzside, MA 00813
    4
                          679 Luna Roads\r\nGreggshire, VA 11650
    396025
                12951 Williams Crossing\r\nJohnnyville, DC 30723
    396026
               0114 Fowler Field Suite 028\r\nRachelborough, ...
    396027
               953 Matthew Points Suite 414\r\nReedfort, NY 7...
               7843 Blake Freeway Apt. 229\r\nNew Michael, FL...
    396028
    396029
                   787 Michelle Causeway\r\nBriannaton, AR 48052
    Name: address, Length: 396030, dtype: object
#Split address into State and Zip code
import re
df[['state', 'zipcode']] = df['address'].str.extract(r'([A-Z]{2}) (\d{5})')
df.drop(['address'], axis=1, inplace=True)
df['state'].nunique() , df['zipcode'].nunique()
\rightarrow \overline{} (54, 10)
df['state'].isna().sum() , df['zipcode'].isna().sum()
```

```
→ (0, 0)
```

```
df['emp_length_yrs'] = df['emp_length'].str.extract('(\d+)')
df.drop(['emp_length'], axis=1, inplace=True)
```

df['term'] = df['term'].str.split().str[0].astype('object')

df.sample()

<b>→</b>		loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	home
	43629	16425.0	36	17.57	590.27	D	D2	Insurance Consultant	

df.shape

# List of numerical columns
num\_cols = df.select\_dtypes(exclude='object')

cat\_cols.sample(3)

<b>→</b>		term	grade	sub_grade	emp_title	home_ownership	verification_statu
	297504	60	Е	E3	Admin Supervisor	RENT	Verifie
	348683	60	С	C3	BUSINESS CONSULTANT	MORTGAGE	Verifie
	361643	60	G	G2	Supply Chain Manager	RENT	Source Verifie

num\_cols.sample(3)

$\overline{\Rightarrow}$		loan_amnt	int_rate	installment	annual_inc	dti	open_acc	pub_rec ı
	225076	14400.0	14.09	492.79	50000.0	34.70	13.0	0
	313487	16000.0	16.55	393.79	70000.0	27.93	6.0	0
								-

num\_cols.skew()

₹	loan_amnt int_rate installment annual_inc dti open_acc pub_rec revol_bal revol_util	0.777285 0.420669 0.983598 41.042725 431.051225 1.213019 6.812303 11.727515 -0.074238
	<b>—</b>	
	total_acc	0.864328
	<pre>mort_acc pub_rec_bankruptcies dtype: float64</pre>	0.412225 12.936099

## 🗸 💡 Insights

· Features are Right skewed

#### Action

• Need to apply log transformations in order to normalise them

```
df1 = df.copy()
df1.sample()
```

<b>→</b>		loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	hon
	241830	8000.0	36	9.99	258.1	В	ВЗ	No Employee Title	

# ♥Q1. What percentage of customers have fully paid their Loan Amount?

df['loan\_status'].value\_counts(normalize=True)\*100

loan\_status
Fully Paid 80.387092
Charged Off 19.612908
Name: proportion, dtype: float64

## 🗸 💡 Insights:

• Target variable distribution is 80%-20%. Data is significantly imbalanced

 $\rightarrow$ 

# Graphical Analysis:


```
uni / bi / multi variate Analysis
```

```
cp = ['indigo','m','darkviolet','magenta','mediumorchid','violet','purple','orchi
num_cols.iloc[:,[0,2,3,4,5,6,8,9,10]].sample()
```

loan\_amnt installment annual\_inc dti open\_acc pub\_rec revol\_util

```
plt.style.use('default')
plt.style.use('seaborn-bright')
outlier_graphical_cols = num_cols.iloc[:,[0,2,3,4,5,6,8,9,10]]
for _,col in enumerate(outlier_graphical_cols.columns):
    plt.figure(figsize=(18,4))
    plt.suptitle(f'plot of {col}',fontsize=15,fontfamily='serif',fontweight='bold plt.subplot(121)
    sns.boxplot(x=df[col],color=cp[_])
    plt.title(f'Boxplot of {col}',fontsize=12,fontfamily='serif',fontweight='bold plt.subplot(122)
    sns.histplot(x=df[col], kde=True,color=cp[_])
    plt.title(f'Distplot of {col}',fontsize=12,fontfamily='serif',fontweight='bol sns.despine()
    plt.show()
```





## ✓ √ Insights:

- 1. The analysis suggests a prevalence of outliers, prompting further investigation into outlier detection techniques.
- 2. Among the numerical features, Potential outliers may still be present.
- 3. Notably, features such as Pub\_rec, Mort\_acc, and Pub\_rec\_bankruptcies display a sparse distribution of unique values, indicating the potential benefit of generating binary features from these variables.

```
#Countplots of various categorical features w.r.t. to target variable loan_status
plt.figure(figsize=(16,17))
plt.suptitle('Countplots of various categorical features w.r.t. to target variabl
             fontsize=14, fontfamily='serif', fontweight='bold', backgroundcolor=cp[
plt.subplot(321)
sns.countplot(data=df, x='loan_status',palette=cp)
plt.title('Loan Status Counts',fontsize=12,fontfamily='serif',fontweight='bold',b
plt.subplot(322)
sns.countplot(data=df, x='loan status', hue='term',palette=cp)
plt.title('Term wise loan status count',fontsize=12,fontfamily='serif',fontweight
plt.subplot(323)
sns.countplot(data=df, x='home_ownership', hue='loan_status',palette=cp)
plt.title('Loan Status Vs Home Ownership',fontsize=12,fontfamily='serif',fontweig
plt.subplot(324)
sns.countplot(data=df, x='verification_status', hue='loan_status',palette=cp)
plt.title('Loan Status Vs Verification Status',fontsize=12,fontfamily='serif',fon
plt.subplot(325)
sns.countplot(data=df, x='issue_month', hue='loan_status',palette=cp)
plt.title('Loan Status Vs issue_month',fontsize=12,fontfamily='serif',fontweight=
plt.subplot(326)
sns.countplot(data=df, x='zipcode', hue='loan_status',palette=cp)
plt.title('Loan Status Vs zipcode',fontsize=12,fontfamily='serif',fontweight='bol
sns.despine()
plt.show()
```





```
zip_codes = ["11650", "86630", "93700"]
states = df[df['zipcode'].isin(zip_codes)]['state']
for zip_code, state in zip(zip_codes, states):
    print(f"Zip code: {zip_code}, State: {state}")

Zip code: 11650, State: VA
    Zip code: 86630, State: MI
    Zip code: 93700, State: MD
```


#### ✓ Observations:


- It's been observed that loans haven't been completely repaid in zip codes 11650, 86630, and 93700.
- Loans haven't been repaid by borrowers residing in 'VA', 'MI', and 'MD'.

```
#Boxplot of various cont. features w.r.t. target variable loan_status
plt.figure(figsize=(18,10))
plt.suptitle('Boxplot of various cont. features w.r.t. target variable loan_statu
             fontsize=14, fontfamily='serif', fontweight='bold', backgroundcolor=cp[
plt.subplot(221)
sns.boxplot(data=df, x='loan_status', y='loan_amnt',palette=cp)
plt.title('Loan Status Vs Loan amounts',fontsize=12,fontfamily='serif',fontweight
plt.subplot(222)
sns.boxplot(data=df, x='loan_status', y='int_rate',palette=cp)
plt.title('Loan Status Vs Interest Rate ',fontsize=12,fontfamily='serif',fontweig
plt.subplot(223)
sns.boxplot(data=df, x='loan_status', y='installment',palette=cp)
plt.title('Loan Status Vs EMI', fontsize=12, fontfamily='serif', fontweight='bold', b
plt.subplot(224)
sns.boxplot(data=df, x='loan_status', y='annual_inc',palette=cp)
plt.ylim(bottom=-5000, top=300000)
plt.title('Loan Status Vs Annual Income',fontsize=12,fontfamily='serif',fontweigh
sns.despine()
plt.show()
```



#### Boxplot of various cont. features w.r.t. target variable loan\_status






- Charged Off customers exhibit a notably higher median interest rate compared to Fully Paid customers.
- The median annual income of Charged Off customers is lower than that of Fully Paid customers.
- Charged Off customers tend to have a higher median EMI compared to Fully Paid customers.
- The median loan amount for Charged Off customers surpasses that of Fully Paid customers.

df.sample()

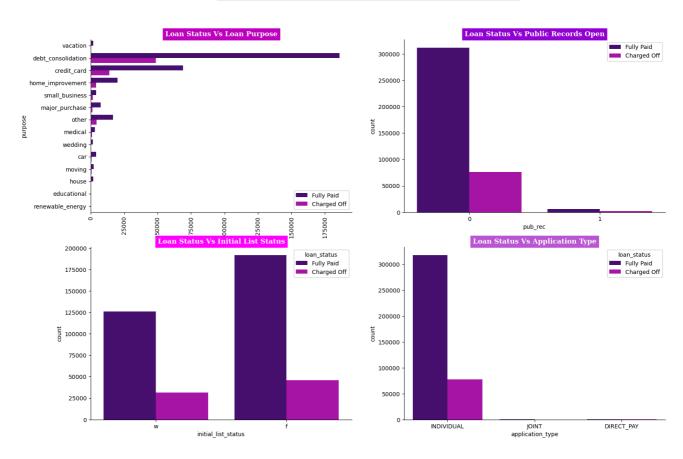
plt.show()

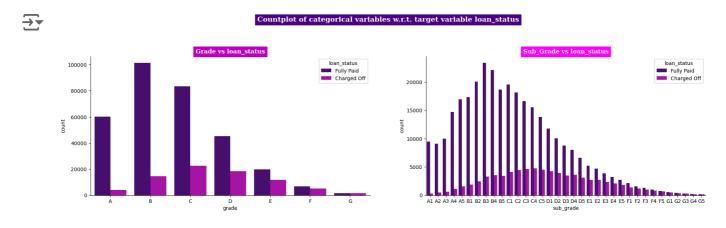


```
#Countplot of categorical variables w.r.t. target variable loan_status
plt.figure(figsize=(18,12))
plt.suptitle('Countplot of categorical variables w.r.t. target variable loan_stat
             fontsize=14, fontfamily='serif', fontweight='bold', backgroundcolor=cp[
plt.subplot(221)
sns.countplot(data=df, y='purpose', hue='loan_status',palette=cp)
plt.xticks(rotation=90)
plt.title('Loan Status Vs Loan Purpose',fontsize=12,fontfamily='serif',fontweight
plt.legend(loc=4)
plt.subplot(222)
sns.countplot(data=df, x='pub_rec',hue='loan_status',palette=cp)
plt.title('Loan Status Vs Public Records Open',fontsize=12,fontfamily='serif',fon
plt.legend(loc=1)
plt.subplot(223)
sns.countplot(data=df, x='initial_list_status', hue='loan_status',palette=cp)
plt.title('Loan Status Vs Initial List Status',fontsize=12,fontfamily='serif',fon
plt.subplot(224)
sns.countplot(data=df, x='application_type',hue='loan_status',palette=cp)
plt.title('Loan Status Vs Application Type',fontsize=12,fontfamily='serif',fontwe
sns.despine()
```



#### $Countplot\ of\ categorical\ variables\ w.r.t.\ target\ variable\ loan\_status$





#### ✓ Observations:

Top 2 loan purpose categories are Debit Consolidation and Credit Card

- Topmost loan type application is INDIVIDUAL
- The distribution of open\_acc appears to be relatively normal when visualized graphically.
- Charged Off and Fully Paid categories exhibit similar distributions.


#### df.sample()



sns.despine()

plt.show()





# Q2. Comment about the correlation between Loan Amount and Installment features.


df[['loan\_amnt', 'installment']].corr()



•		loan_amnt	installment	
	loan_amnt	1.000000	0.953929	
	installment	0.953929	1.000000	

```
plt.figure(figsize = (15,5))
sns.scatterplot(data = df, x = 'loan_amnt', y = 'installment', alpha = 0.5, hue =
plt.title('Loan Amt Vs Installments',fontsize=12,fontfamily='serif',fontweight='b
sns.despine()
plt.show()
```





## Insights:

The correlation coefficient measures the strength and direction of the linear relationship between two variables. In this case, the correlation coefficient between 'loan\_amnt' and 'installment' is quite high, approximately 0.95, indicating a strong positive linear relationship between these two variables.

- Loan Terms: Understanding the relationship between loan amount and installment
  payments is crucial for setting appropriate loan terms. Lenders can adjust loan terms such
  as interest rates and repayment periods based on the borrower's ability to handle
  installment payments associated with different loan amounts.
- Potential Multicollinearity: When building predictive models, it's essential to be cautious of
  multicollinearity between highly correlated predictor variables. Multicollinearity can lead to
  unstable estimates and difficulties in interpreting the model coefficients. Therefore, it
  might be necessary to address multicollinearity through techniques such as variable
  selection or regularization.
- Q3. The majority of people have home ownership as \_\_\_\_.

(df['home\_ownership'].value\_counts(normalize=True)\*100).to\_frame()



#### proportion

home	owne	rship

MORTGAGE	50.084085
RENT	40.347953
OWN	9.531096
OTHER	0.028281
NONE	0.007828
ANY	0.000758

## Insights:

- Mortgage holders comprise the majority with approximately 50.08%, indicating that a significant portion of individuals own homes through Mortgage agreements.
- Renters constitute a substantial portion, accounting for around 40.35% of home ownership types. This suggests a sizable demographic of individuals who opt for renting rather than owning a home.

Q4. People with grades 'A' are more likely to fully pay their loan. (T/F)

pd.crosstab(df['grade'],df['loan\_status'], normalize = 'index')

<b>→</b>	loan_status grade	Charged Off	Fully Paid
	Α	0.062879	0.937121
	В	0.125730	0.874270
	С	0.211809	0.788191
	D	0.288678	0.711322

**F** 0.427880 0.572120 **G** 0.478389 0.521611

0.373634

## Insights:

Ε

0.626366

- True . Grade 'A' borrowers demonstrate a significantly high likelihood of fully repaying their loans, with approximately 93.71% of loans being fully paid. This suggests that borrowers with the highest credit rating are more inclined to fulfill their loan obligations successfully.
- The proportion of charged-off loans for grade 'A' borrowers is relatively low, standing at approximately 6.29%. This indicates a low default rate among borrowers with the highest credit rating, emphasizing their creditworthiness and reliability in loan repayment.
- √ Q5. Name the top 2 afforded job titles.

df[df['emp\_title'] != 'No Employee Title']['emp\_title'].value\_counts().to\_frame()

<b>→</b>		count
	emp_title	
	Teacher	4389
	Manager	4250
	Registered Nurse	1856
	RN	1846
	Supervisor	1830


df.groupby('emp\_title')['loan\_status'].count().sort\_values(ascending=False).to\_fr

<b>→</b>		loan_status
	emp_title	
	Teacher	4389
	Manager	4250
	Registered Nurse	1856
	RN	1846
	Supervisor	1830

- - The Most afforded job titles are Teachers & Managers.

```
plt.figure(figsize=(20,12))
sns.heatmap(num_cols.corr(), annot=True, cmap='Purples')
plt.title('Correlations - Heatmap',fontsize=12,fontfamily='serif',fontweight='bol
plt.show()
```





#### 

- There exists a strong correlation between loan\_amnt and installment, indicating that higher loan amounts correspond to larger installment payments.
- The variables total\_acc and open\_acc exhibit a significant correlation.
- There is a notable correlation between pub\_rec\_bankruptcies and pub\_rec.

#### Qutlier Treatment:

```
numerical_cols = df.select_dtypes(include=np.number).columns
numerical_cols
    Index(['loan_amnt', 'int_rate', 'installment', 'annual_inc', 'dti',
    'open_acc',
            'pub_rec', 'revol_bal', 'revol_util', 'total_acc', 'mort_acc',
            'pub_rec_bankruptcies'],
          dtype='object')
# outlier treatment
def remove_outliers_zscore(df, threshold=2): #(considering 2 std.dev away from me
   Remove outliers from a DataFrame using the Z-score method.
    Parameters:
        df (DataFrame): The input DataFrame.
        threshold (float): The Z-score threshold for identifying outliers.
                           Observations with a Z-score greater than this threshol
                           will be considered as outliers.
   Returns:
       DataFrame: The DataFrame with outliers removed.
   # Calculate Z-scores for numerical columns
   z_scores = (df[numerical_cols] - df[numerical_cols].mean()) / df[numerical_co
   # Identify outliers
   outliers = np.abs(z_scores) > threshold
   # Keep non-outliers for numerical columns
   df_cleaned = df[~outliers.any(axis=1)]
    return df_cleaned
cleaned_df = remove_outliers_zscore(df1)
print(cleaned_df.shape)
```

```
(311392, 30)
\rightarrow
def clip_outliers_zscore(df, threshold=2):
    Clip outliers in a DataFrame using the Z-score method.
    Parameters:
        df (DataFrame): The input DataFrame.
        threshold (float): The Z-score threshold for identifying outliers.
                           Observations with a Z-score greater than this threshol
                           will be considered as outliers.
    Returns:
        DataFrame: The DataFrame with outliers clipped.
   # Calculate Z-scores for numerical columns
    z_scores = (df[numerical_cols] - df[numerical_cols].mean()) / df[numerical_co
   # Clip outliers
    clipped_values = df[numerical_cols].clip(df[numerical_cols].mean() - threshol
                                              df[numerical_cols].mean() + threshol
                                              axis=1)
   # Assign clipped values to original DataFrame
    df clipped = df.copy()
    df_clipped[numerical_cols] = clipped_values
    return df clipped
clipped_df = clip_outliers_zscore(df1)
print(clipped_df.shape)
    (396030, 30)
data = cleaned_df.copy()
cp_data = clipped_df.copy()
data.sample()
\rightarrow
            loan_amnt term int_rate installment grade sub_grade emp_title hom
     110850
                14000.0
                          36
                                  11.67
                                               462.8
                                                         В
                                                                   B4
                                                                         Manager
data['pub_rec_bankruptcies'].value_counts() , data['pub_rec'].value_counts()
    (pub_rec_bankruptcies
          311392
     Name: count, dtype: int64,
     pub rec
          311392
     Name: count, dtype: int64)
```

cp\_data['pub\_rec\_bankruptcies'].value\_counts() , cp\_data['pub\_rec'].value\_counts()

data.shape

→ (311392, 30)

data.info()

<<class 'pandas.core.frame.DataFrame'>
 Index: 311392 entries, 0 to 396029
 Data columns (total 30 columns):

#	Column	Non-Nu	ll Count	Dtype
0	loan_amnt	311392	non-null	float64
1	term	311392	non-null	object
2	int_rate	311392	non-null	float64
3	installment	311392	non-null	float64
4	grade	311392	non-null	object
5	sub_grade	311392	non-null	object
6	emp_title	311392	non-null	object
7	home_ownership	311392	non-null	object
8	annual_inc	311392	non-null	float64
9	verification_status	311392	non-null	object
10	loan_status	311392	non-null	object
11	purpose	311392	non-null	object
12	title	311392	non-null	object
13	dti	311392		float64
14	open_acc	311392	non-null	float64
15	pub_rec	311392	non-null	int64
16	revol_bal	311392		float64
17	revol_util	311392		float64
18	total_acc	311392		float64
19	initial_list_status	311392	non-null	object
20	application_type	311392	non-null	object
21	mort_acc	311392	non-null	int64
22	<pre>pub_rec_bankruptcies</pre>		non-null	int64
23	issue_month		non-null	object
24	issue_year		non-null	object
25	er_cr_line_m	311392	non-null	object
26	er_cr_line_y	311392	non-null	object
27	state		non-null	object
28	zipcode		non-null	object
29	emp_length_yrs		non-null	object
		3) <b>,</b> obje	ect(18)	
memo	ry usage: 73.6+ MB			

### Manual encoding:

```
data['loan_status']=data.loan_status.map({'Fully Paid':1, 'Charged Off':0})
data['initial_list_status']=data.initial_list_status.map({'w':0, 'f':1})
data.head()
```

<b>→</b>		loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	home_ov
	0	10000.0	36	11.44	329.48	В	B4	Marketing	
	1	8000.0	36	11.99	265.68	В	B5	Credit analyst	МС
	2	15600.0	36	10.49	506.97	В	ВЗ	Statistician	
	3	7200.0	36	6.49	220.65	А	A2	Client Advocate	
	4	24375.0	60	17.27	609.33	С	C5	Destiny Management Inc.	MC

- ✓ Feature selection done by hypothesis testing & VIF(multicolinearity)
  - Find VIF after modelling and remove features with high VIF (>5):

```
def calc_vif(X):
    # Calculating the VIF
    vif=pd.DataFrame()
    vif['Feature']=X.columns
    vif['VIF']=[variance_inflation_factor(X.values,i) for i in range(X.shape[1])]
    vif['VIF']=round(vif['VIF'],2)
    vif=vif.sort_values(by='VIF',ascending=False)
    return vif
cat_cols = data.select_dtypes(include=['object']).columns.tolist()
for col in cat_cols:
   chi2, p, dof, expected = chi2_contingency(pd.crosstab(data[col], data['loan_s
    if p > 0.05:
        print('>>>>> Independent feature - Not Significant:',col,' >> p value:'
>>>>> Independent feature - Not Significant: emp_title >> p value: 0.536717
    >>>>> Independent feature - Not Significant: title >> p value: 1.0
    >>>>> Independent feature - Not Significant: er_cr_line_m >> p value: 0.27
    >>>>> Independent feature - Not Significant: state >> p value: 0.760478089
```

**→** (311392, 19)

lt.sample()



#### Performing OneHotEncoding on feature having multiple variable
dummies=['zipcode', 'grade', 'purpose', 'home\_ownership', 'verification\_status', 'app
ltd = pd.get\_dummies(lt, columns=dummies, drop\_first=True)\*1

ltd.shape

→ (311392, 50)

ltd.dtypes

_	
loan_amnt	float64
term	object
int_rate	float64
installment	float64
annual_inc	float64
loan_status	int64
dti	float64
open_acc	float64
revol_bal	float64
revol_util	float64
total_acc	float64
mort_acc	int64
emp_length_yrs	object
zipcode_05113	int64
zipcode_11650	int64
zipcode_22690	int64
zipcode_29597	int64
	int64
zipcode_48052	int64
zipcode_70466	int64
zipcode_86630	int64
zipcode_93700	int64
grade_B	int64
grade_C	int64
grade_D	int64
grade_E	int64
grade_F	int64
grade_G	int64
purpose_credit_card	int64
purpose_debt_consolidation	int64
	<pre>int_rate installment annual_inc loan_status dti open_acc revol_bal revol_util total_acc mort_acc emp_length_yrs zipcode_05113 zipcode_11650 zipcode_22690 zipcode_29597 zipcode_30723 zipcode_48052 zipcode_48052 zipcode_70466 zipcode_86630 zipcode_93700 grade_B grade_C grade_D grade_E grade_F grade_G purpose_credit_card</pre>

purpose_educational	int64
<pre>purpose_home_improvement</pre>	int64
purpose_house	int64
<pre>purpose_major_purchase</pre>	int64
purpose_medical	int64
purpose_moving	int64
purpose_other	int64
purpose_renewable_energy	int64
purpose_small_business	int64
purpose_vacation	int64
purpose_wedding	int64
home_ownership_MORTGAGE	int64
home_ownership_NONE	int64
home_ownership_OTHER	int64
home_ownership_OWN	int64
home_ownership_RENT	int64
verification_status_Source Verified	int64
verification_status_Verified	int64
application_type_INDIVIDUAL	int64
application_type_JOINT	int64
dtype: object	

ltd.sample(8)



	loan_amnt	term	int_rate	installment	annual_inc	loan_status	dti	(
118504	15000.0	36	10.99	491.01	85000.0	0	17.05	-
20036	26000.0	36	11.99	863.45	100000.0	1	13.22	
388815	9175.0	36	13.35	310.70	40000.0	1	15.12	
388094	13000.0	60	18.24	331.82	82000.0	1	18.29	
254903	9600.0	36	15.31	334.25	65000.0	1	11.78	
264585	22500.0	36	15.61	786.71	81000.0	1	18.42	
368842	15000.0	36	8.90	476.30	120000.0	1	9.56	
44417	24000.0	60	13.99	558.32	75000.0	1	16.26	

### ✓ Model:

#Prepare X and y dataset i.e. independent and dependent datasets

```
X = ltd.drop(['loan_status'], axis=1)
```

y = ltd['loan\_status']

```
#Split the data into train and test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)

$\frac{(249113, 49)}{(62279, 49)}(249113,)(62279,)}$
```

### Minmax scaling the data

```
scaler = MinMaxScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
X_train = pd.DataFrame(X_train, columns=X.columns)
X_test = pd.DataFrame(X_test, columns=X.columns)
```

#### X\_train.head()

 $\overline{\mathbf{T}}$ 

•		loan_amnt	term	int_rate	installment	annual_inc	dti	open_acc	revol <sub>.</sub>
	0	0.379538	0.0	0.339161	0.411590	0.207250	0.465341	0.368421	0.17
	1	0.643564	1.0	0.680070	0.524221	0.367868	0.252652	0.473684	0.22
	2	0.168317	0.0	0.208625	0.176198	0.134712	0.357576	0.368421	0.05
	3	0.379538	1.0	0.680070	0.307444	0.367868	0.449242	0.315789	0.25
	4	0.368812	0.0	0.543706	0.421460	0.246109	0.315530	0.263158	0.09

## ✓ Model-1

```
#Fit the Model on training data
logreg_model = LogisticRegression()
logreg_model.fit(X_train, y_train)
```

```
▼ LogisticRegression ()
```

```
#Predit the data on test dataset
y_train_pred = logreg_model.predict(X_train)
y_test_pred = logreg_model.predict(X_test)
```

If logreg\_model.score(X\_test, y\_test) consistently returns 1, it would imply that your model is predicting the test set perfectly, which could be a sign of overfitting, data leakage, or an issue with the evaluation process.

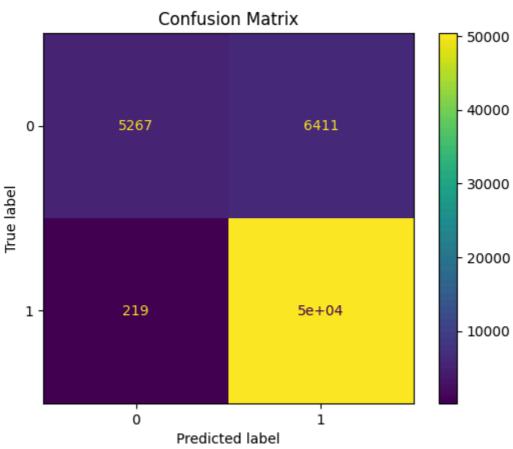
```
#Model Evaluation
print('Train Accuracy :', logreg_model.score(X_train, y_train).round(2))
print('Train F1 Score:',f1_score(y_train,y_train_pred).round(2))
print('Train Recall Score:',recall_score(y_train,y_train_pred).round(2))
print('Train Precision Score:',precision_score(y_train,y_train_pred).round(2))

print('NTest Accuracy :',logreg_model.score(X_test,y_test).round(2))
print('Test F1 Score:',f1_score(y_test,y_test_pred).round(2))
print('Test Recall Score:',recall_score(y_test,y_test_pred).round(2))
print('Test Precision Score:',precision_score(y_test,y_test_pred).round(2))

# Confusion Matrix
cm = confusion_matrix(y_test, y_test_pred)
disp = ConfusionMatrixDisplay(cm)
disp.plot()
plt.title('Confusion Matrix')
plt.show()
```

Train Accuracy: 0.89
Train F1 Score: 0.94
Train Recall Score: 1.0
Train Precision Score: 0.89

Test Accuracy: 0.89
Test F1 Score: 0.94
Test Recall Score: 1.0
Test Precision Score: 0.89



print(classification\_report(y\_test,y\_test\_pred))

<b>→</b>		precision	recall	f1-score	support
	0 1	0.96 0.89	0.45 1.00	0.61 0.94	11678 50601
	accuracy macro avg weighted avg	0.92 0.90	0.72 0.89	0.89 0.78 0.88	62279 62279 62279

• Here the recall value for the 'charged off' is very low, Hence will build a better model


#### ✓ Model-2

# Oversampling to balance the target variable sm=SMOTE(random state=42) X\_train\_res, y\_train\_res = sm.fit\_resample(X\_train,y\_train.ravel()) print(f"Before OverSampling, count of label 1: {sum(y train == 1)}") print(f"Before OverSampling, count of label 0: {sum(y train == 0)}") print(f"After OverSampling, count of label 1: {sum(y\_train\_res == 1)}") print(f"After OverSampling, count of label 0: {sum(y train res == 0)}") Before OverSampling, count of label 1: 202401 Before OverSampling, count of label 0: 46712 After OverSampling, count of label 1: 202401 After OverSampling, count of label 0: 202401 model = LogisticRegression() model.fit(X\_train\_res, y\_train\_res) train preds = model.predict(X train) test\_preds = model.predict(X\_test) #Model Evaluation print('Train Accuracy :', model.score(X\_train, y\_train).round(2)) print('Train F1 Score:',f1\_score(y\_train,train\_preds).round(2)) print('Train Recall Score:',recall\_score(y\_train,train\_preds).round(2)) print('Train Precision Score:',precision\_score(y\_train,train\_preds).round(2)) print('\nTest Accuracy :',model.score(X\_test,y\_test).round(2)) print('Test F1 Score:',f1 score(y test,test preds).round(2)) print('Test Recall Score:',recall\_score(y\_test,test\_preds).round(2)) print('Test Precision Score:',precision\_score(y\_test,test\_preds).round(2)) # Confusion Matrix cm = confusion\_matrix(y\_test, test\_preds) disp = ConfusionMatrixDisplay(cm) disp.plot() plt.title('Confusion Matrix') plt.show()

Train Accuracy: 0.79
Train F1 Score: 0.86

Train Recall Score: 0.79
Train Precision Score: 0.95

Test Accuracy: 0.8 Test F1 Score: 0.86 Test Recall Score: 0.79 Test Precision Score: 0.95



y\_pred = test\_preds
print(classification\_report(y\_test,y\_pred))

<b>→</b>		precision	recall	f1-score	support
	0 1	0.47 0.95	0.81 0.79	0.60 0.86	11678 50601
	accuracy macro avg weighted avg	0.71 0.86	0.80 0.80	0.80 0.73 0.81	62279 62279 62279

## Observations:

 The model demonstrates a high recall score, successfully identifying 80% of actual defaulters.

- However, the precision for the positive class (defaulters) is low; only 47% of predicted defaulters are actually defaulters.
- This high recall and low precision indicate that while the model is effective at flagging
  most defaulters, it also results in many false positives. Consequently, many deserving
  customers may be denied loans.
- The low precision adversely affects the F1 score, reducing it to 60%, despite an overall accuracy of 80%. This highlights the trade-off between precision and recall in the model's performance.

### Explanation:

- The model is good at catching most people who don't pay back their loans it catches 80% of them.
- But, when it says someone won't pay back, it's right only half of the time.47% So, there's a chance it's making mistakes and wrongly flagging people.
- Because of these mistakes, some people who deserve loans might not get them.
- Even though the model seems okay overall, its balance between being right and not making mistakes isn't great. It's like a seesaw; when one side goes up, the other goes down.

## Regularization Model

test\_scores.append(te\_score)

```
#Try with different regularization factor lamda and choose the best to build the
lamb = np.arange(0.01, 10000, 10)

train_scores = []

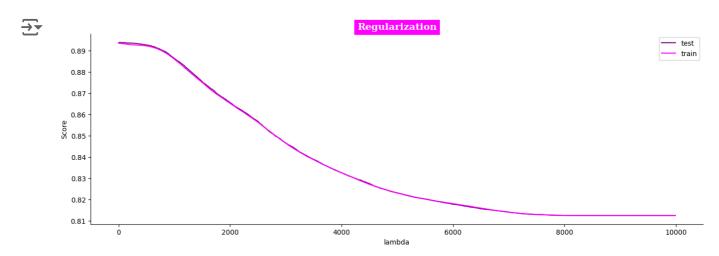
test_scores = []

for lam in lamb:
    model = LogisticRegression(C = 1/lam)
    model.fit(X_train, y_train)

    tr_score = model.score(X_train, y_train)
    te_score = model.score(X_test, y_test)

    train_scores.append(tr_score)
```

```
#Plot the train and test scores with respect lambda values i.e. regularization fa
ran = np.arange(0.01, 10000, 10)
plt.figure(figsize=(16,5))
sns.lineplot(x=ran,y=test_scores,color='purple',label='test')
sns.lineplot(x=ran,y=train_scores,color='magenta',label='train')
plt.title('Regularization',fontsize=14,fontfamily='serif',fontweight='bold',backg
plt.xlabel("lambda")
plt.ylabel("Score")
sns.despine()
plt.show()
```



#Check the index of best test score and the check the best test score

```
print(np.argmax(test_scores))
print(test_scores[np.argmax(test_scores)])
```

2 0.8939289327060486

#Calculate the best lambda value based on the index of best test score

best\_lamb = 
$$0.01 + (10*2)$$
  
best\_lamb

<del>\_</del> 20.01

```
#Fit the model using best lambda
```

```
reg_model = LogisticRegression(C=1/best_lamb)
reg_model.fit(X_train, y_train)
```

 $\rightarrow$ 

```
LogisticRegression
LogisticRegression(C=0.04997501249375312)
```

#Predict the y\_values and y\_probability values

```
y_reg_pred = reg_model.predict(X_test)
y_reg_pred_proba = reg_model.predict_proba(X_test)
```

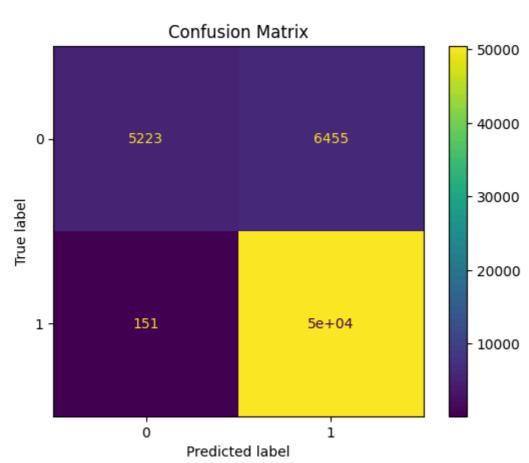
#Print model score

print(f'Logistic Regression Model Score with best lambda: ',end='')
print(round(model.score(X\_test, y\_test)\*100,2),'%')

Logistic Regression Model Score with best lambda: 89.39 %

```
# Confusion Matrix
cm = confusion_matrix(y_test, y_reg_pred)
disp = ConfusionMatrixDisplay(cm)
disp.plot()
plt.title('Confusion Matrix')
plt.show()
```





print(classification\_report(y\_test, y\_reg\_pred))

<b>→</b>	precision	recall	f1-score	support
0 1	0.97 0.89	0.45 1.00	0.61 0.94	11678 50601
accuracy macro avg weighted avg	0.93 0.90	0.72 0.89	0.89 0.78 0.88	62279 62279 62279

Observations from classification report:

### Regularized model

• Precision: 89% • Recall: 100% • F1-score: 94% Accuracy: 89%

# K-fold - Cross\_validation

cross validation accuracy has to be approx 89%

```
x=scaler.fit_transform(X)
kfold = KFold(n_splits=10)
accuracy = np.mean(cross_val_score(reg_model,x,y,cv=kfold,scoring='accuracy'))
print("Cross Validation accuracy : {:.3f}".format(accuracy))
Cross Validation accuracy: 0.894
cm = confusion_matrix(y_test, y_reg_pred)
cm_df = pd.DataFrame(cm, index=['Defaulter','Fully paid'], columns=['Defaulter','
cm_df
\rightarrow
               Defaulter Fully paid
     Defaulter
                    5223
                                 6455
     Fully paid
```

## 

- TN = 5223 (True Negative: Correctly predicted Charged Off)
- TP = 50450 (True Positive: Correctly predicted Fully Paid)

151

FP = 6455 (False Positive: Predicted Fully Paid but actually Charged Off)

50450

- FN = 151 (False Negative: Predicted Charged Off but actually Fully Paid)
- Actual Negative (Charged Off) = 5223 + 6455 = 11678
- Actual Positive (Fully Paid) = 151 + 50450 = 50601
- Predicted Negative (Charged Off) = 5223 + 151 = 5374
- Predicted Positive (Fully Paid) = 6455 + 50450 = 56905

```
#Collect the model coefficients and print those in dataframe format
coeff_df = pd.DataFrame()
coeff_df['Features'] = X_train_res.columns
coeff_df['Weights'] = model.coef_[0]
coeff_df['ABS_Weights'] = abs(coeff_df['Weights'])
coeff_df = coeff_df.sort_values(['ABS_Weights'], ascending=False)
coeff_df
```